

Final Project Specification
April 17, 2015

Project Title: “Machine Learning for Tic-Tac-Toe”

Team Members: Vincent Chow (chow@college.harvard.edu), Stephen Albro (salbro@college.harvard.edu), Brandon Price (bprice@college.harvard.edu), and Peter Hickman (phickman@college.harvard.edu)

TF Advisor: Armaghan Behlum

Detailed Description:

Signatures/Interfaces (see our [GitHub Repository](#)) with pseudocode -

- initializer.py
 - types:
 - type player = X | O
 - type action = int * int (* coordinates on gameboard *)
 - type board = (*array of arrays *)
 - type state = board * player
 - functions:
 - new () : state = (*empty board & player to move first*)
 - nextState(state, action) : state
 - Eval (state): player | tie | continue
 - opponent (player) = (* other player *)
 - valid(state, action) : bool = (* checks whether action is legal *)
- matrices.py
 - functions
 - makeKey (state) : string (* takes state/creates string representation *)
 - nextKey (state) : makeKey (nextState, opponent(player))
 - addKey (string) : (* adds row to rewardTable & qTable with this key, reward is 0 and q-value is 0 *)
 - setup matrices
 - type stateTable = dict (* dictionary with states as keys and actTables as values *)
 - type actTable = dict (* dictionary with squares as keys and floats as values *)
 - val mutable qTable : stateTable = (* floats [keys in inner actTable] are q-values; we'll update this *)
 - val rewardTable : stateTable = (* floats [keys in inner actTable] are rewards [0, 1, or -1]; don't update this *)

- value (board, square, qTable): float (* looks up q-value for board and next move square*)
- learning.py
 - chooseMove (state) : (state*square)= (* identifies the state that maximizes the Q* algorithm and makes that move *)
 - val gamma : int = 0.8 (* discount factor *)
 - updateQvalues (stateKey, action, nextKey, reward) = (* updates the qTable using the function $Q^*(s,a) \leftarrow r + \gamma \min_{a'} Q^*(s',a')$ *)
- graphics.py
 - functions
 - draw (state) : (displays gameboard)
- main.py (main file calls other modules and runs machine learning program)
 - gameLearning (maxGames) (* plays games until maxGames reached *)

Timeline:

- Week of 04/12 - 4/18:
 - setup distributed version control
 - create signatures/interfaces for modules
 - understand Q learning
 - begin implementations and pseudocode
- Week of 04/19 - 4/25:
 - drawn board interacts with user event (e.g. mouse clicks)
 - modules partially implemented
 - e.g. progress made on updating Q learning algorithm
 - e.g. capable of instantiating statetables
- Week of 04/26 - 5/2:
 - modules fully implemented
 - program capable of playing two algorithms against each other
 - graphics visualizes the game of two algorithms playing against each other
- Week of 5/3 - 5/6:
 - Humans compete against tic-tac-toe
 - Compare versions with different chooseMove and updateQvalue functions
 - Have different versions of algorithm competing
 - Reach: implement checkers
 - Create new reward matrix
 - Different set of valid moves, new dictionaries

Progress Report:

- We have successfully setup a git repository for collaborating on this project, and given Army access to view our code
- GUI interface has been partially implemented; capable of rendering a state of a tic-tac-toe board.
- Understanding of Q-learning and TD learning has been achieved, and we have pseudocode and signatures to begin implementing these algorithms
- We've found solid reference material from textbooks and other courses to teach us what we know about Reinforcement Learning and its application to Tic-tac-toe.
- We have two ideas for how to implement the chooseMove function, one using modified roulette wheel selection and the other using a (as yet) un-named formula from Tom Mitchell's *Machine Learning*.